

Reduce JavaScript Payloads with Tree Shaking



By Jeremy Wagner

Jeremy is a contributor to WebFundamentals

Today's web applications can get pretty big, especially the JavaScript part of them. As of mid-2018, HTTP Archive puts the median transfer size of JavaScript on mobile devices at approximately 350 KB. And this is just transfer size! JavaScript is often compressed when sent over the network, meaning that the *actual* amount of JavaScript is quite a bit more after the browser decompresses it. That's important to point out, because as far as resource *processing* is concerned, compression is irrelevant. 900 KB of decompressed JavaScript is still 900 KB to the parser and compiler, even though it may be ~300 KB when compressed.



Figure 1. The process of downloading and running JavaScript. Note that even though the transfer size of the script is 300 KB compressed, it is still 900 KB worth of JavaScript that must be parsed, compiled, and executed.

JavaScript is an expensive resource to process. Unlike images which only incur relatively trivial decode time once downloaded, JavaScript must be parsed, compiled, and then finally executed. Byte for byte, this makes JavaScript more expensive than other types of resources.

JavaScript bytes

~170KB

!==

JPEG bytes

~170KB

Network Transmission

main-javascript-bundle.js

3.4 s 170KB

photo.jpg


3.4 s 170KB

Resource Processing


Flame Chart	Bottom-Up	Call Tree	Event Log
Main	Filter	Group by URL	
Self Time	Total Time	Activity	
1997.0 ms 44.6 %	1997.0 ms 44.6 %	native V8Runtime	
608.9 ms 13.6 %	937.6 ms 20.9 %	Compile	
303.6 ms 6.8 %	310.3 ms 6.9 %	Parse	

~2s in Parse/Compile

~1.5s in Execution



0.064s in Image Decode



0.028s in Rasterize Paint

@addyosmani - 170KB of (compressed) JS vs. JPEG bytes over a slow 3G network on a Moto G4. JS needing parsed is even larger once decompressed.

Figure 2. The processing cost of parsing/compiling 170 KB of JavaScript vs decode time of an equivalently sized JPEG. ([source](#)).

While improvements are continually being made to improve the efficiency of JavaScript engines, improving JavaScript performance is, as always, a task fit for developers. After all, who better to improve application architecture than the architects themselves?

To that end, there are techniques to improve JavaScript performance. Code splitting, is one such technique that improves performance by partitioning application JavaScript into chunks, and serving those chunks to only the routes of an application that need them. This technique works, but it doesn't address a common problem of JavaScript-heavy applications, which is the inclusion of code that's never used. To solve this problem, we rely on tree shaking.

What is tree shaking?

Tree shaking is a form of dead code elimination. The term was popularized by Rollup, but the concept of dead code elimination has existed for some time. The concept has also found purchase in webpack, which is demonstrated in this article by way of a sample app.

The term "tree shaking" comes from the mental model of your application and its dependencies as a tree-like structure. Each node in the tree represents a dependency that

provides distinct functionality for your app. In modern apps, these dependencies are brought in via static import statements like so:

```
// Import all the array utilities!  
import arrayUtils from "array-utils";
```



Note: If you're not sure what ES6 modules are, I highly recommend [this excellent explainer over at Pony Foo](#). This guide assumes you have working knowledge of how ES6 modules work, so if you don't know anything about them, give that article a read!

When your app is young (a sapling, if you will), you may have relatively few dependencies. You're also using most (if not all) the dependencies you add. As your app ages, however, more dependencies can get added. To compound matters, older dependencies fall out of use, but may not get pruned from your codebase. The end result is that an app ends up shipping with a lot of unused JavaScript. Tree shaking addresses this by taking advantage of how we use static `import` statements to pull in specific parts of ES6 modules:

```
// Import only some of the utilities!  
import { unique, implode, explode } from "array-utils";
```



The difference between this `import` example and the previous one is that rather than importing *everything* from the "array-utils" module (which could be a lot of stuff!), this example imports only specific parts of it. In dev builds, this doesn't really change anything, as the entire module gets imported regardless. In production builds, however, we can configure webpack to "shake" off exports from ES6 modules that weren't explicitly imported, making those production builds smaller. In this guide, you're going to learn how to do just that!

Finding opportunities to shake a tree

For illustrative purposes, I created [a sample one-page app](#) that uses webpack to demonstrate how tree shaking works. You can clone it and follow along if you like, but we'll cover every step of the way together in this guide, so cloning isn't necessary (unless hands-on learning is your thing).

The sample app is a super simple searchable database of guitar effect pedals. You enter a query and a list of effect pedals pops up.

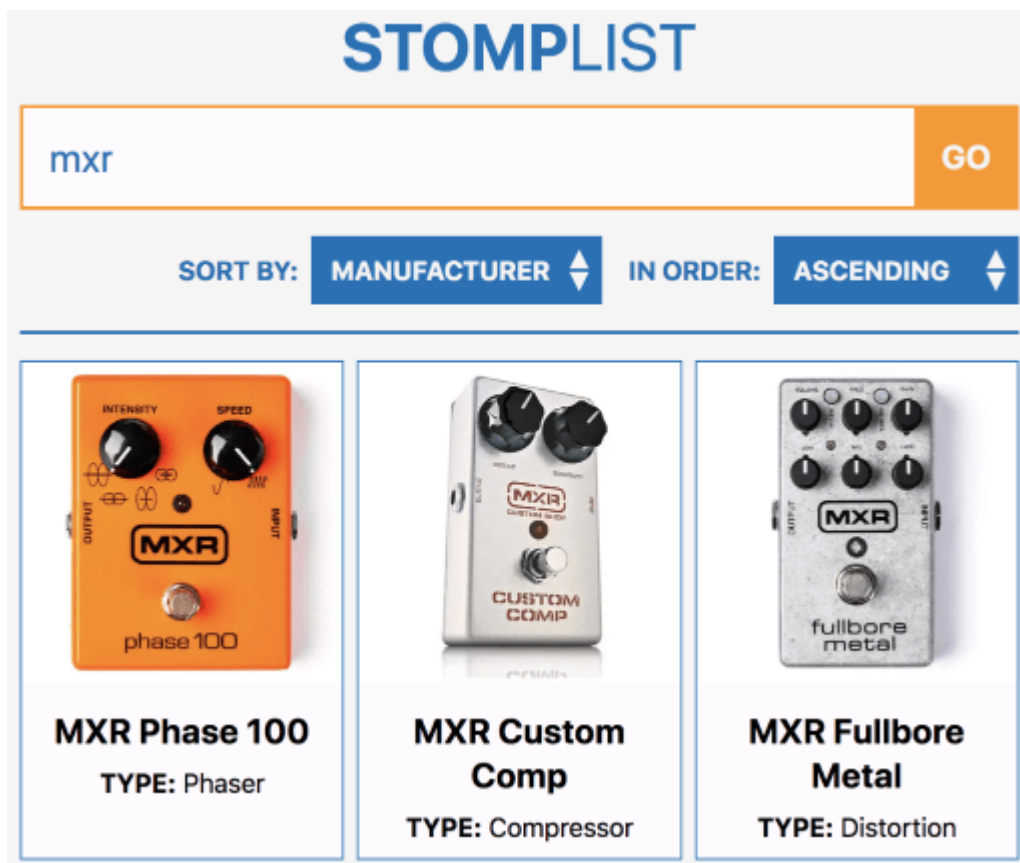


Figure 3. A screenshot of the sample app.

Predictably, the behavior that drives this app is separated into vendor (i.e., Preact and Emotion) and app-specific code bundles (or "chunks", as webpack calls them):



Name	Size
 vendors.6bb454d5.js	37.4 KB
 main.797ebb8b.js	21.5 KB

Figure 4. The app's two JavaScript bundles. These are uncompressed sizes.

The JavaScript bundles shown in the figure above are production builds, meaning they're optimized through uglification. 21.1 KB for an app-specific bundle isn't bad (like *at all*). But! It should be noted that no tree shaking is occurring whatsoever. Let's look at the app code and see what we can do to fix that.

Note: If you don't care for long-winded explanations and just want to dive into code, you can go ahead and check out [the tree-shake branch](#) in the app's GitHub repo. You can also [diff this branch master](#) to see exactly what was changed to make tree shaking work!

In any application, finding opportunities for tree shaking are going to involve looking for static `import` statements. Near the top of the main component file, you'll see a line like this:

```
import * as utils from "../../utils/utils";
```



Maybe you've seen something like this before. The number of ways ES6 module exports can be imported are numerous, but ones like this should get your attention. This specific line is saying "Hey, `import everything` from the `utils` module, and put it in a namespace called `utils`." The big question to ask here is, "just how much *stuff* is in that module?"

Well, if you look at the `utils` module source code, you'll find there's a *lot*. Like around 1,300 lines of code.

Okay, don't worry. Maybe all that stuff is being used, right? *Do we need all that stuff?* Let's double check by searching the main component file we imported the `utils` module into and see how many instances of that namespace come up. Surely, we must be using all that of that stuff for *something*.




Figure 5. The `utils` namespace we've imported tons of modules from is only invoked three times within the main component file.

Well, *that's* no good. We're only using the `utils` namespace in three spots in our application code. But for what functions? If we take a look at the main component file again, it appears to be only one function, which is `utils.simpleSort`, which is used to sort the search results list by a number of criteria when sorting dropdowns are changed:

```
if (this.state.sortBy === "model") {  
  // Simple sort gets used here...  
  json = utils.simpleSort(json, "model", this.state.sortOrder);  
} else if (this.state.sortBy === "type") {  
  // ..and here...  
  json = utils.simpleSort(json, "type", this.state.sortOrder);  
} else {  
  // ..and here.  
  json = utils.simpleSort(json, "manufacturer", this.state.sortOrder);  
}
```



So, that's just great. Out of a 1,300 line file with a bunch of exports, I'm only using one of them. Turns out I'm pretty bad at webperfs.

Note: This project is purposefully kept simple, so it's pretty easy in this case to find out where the bloat is coming from. In large projects with many modules, however, it's tough to find out how much of a bundle is made up of which imports. Tools such as [Webpack Bundle Analyzer](#) and [source-map-explorer](#) can help, but assistive tooling is still being developed to fill this need.

Of course, now would be the time to admit this example is a *bit* manufactured for this article's benefit. While that's absolutely the case here, it doesn't change the fact that this synthetic sort of scenario resembles actual optimization opportunities you may encounter in your own very real apps. So now that you've identified an opportunity for tree shaking to be useful, how do we actually *do* it?

Keeping Babel from transpiling ES6 modules to CommonJS modules

Babel is an indispensable tool most apps need. Unfortunately, it can also make straightforward tasks like tree shaking a bit more difficult, precisely *because* of what it does for us. If you're using babel-preset-env, one thing it automatically does for you is transpile your nice ES6 modules into more widely compatible CommonJS modules (i.e., modules you `require` instead of `import`). This is great and all until we want to start tree shaking.

The problem is that tree shaking is much more difficult to do for CommonJS modules, and webpack won't know what to prune from the bundle if you decide to use them. The solution is simple: We configure `babel-preset-env` to leave ES6 modules alone. Wherever you configure Babel (be it in `.babelrc` or `package.json`) this means adding a little something extra:

```
{
  "presets": [
    ["env", {
      "modules": false
    }]
  ]
}
```



Simply specifying `"modules": false` in your `babel-preset-env` config gets Babel to behave how we want, which allows webpack to analyze your dependency tree and shake off those unused dependencies. Furthermore, this process doesn't cause compatibility issues, as webpack ends up converting your code into a format that's widely compatible, anyway.

Keeping side effects in mind

Another aspect to consider when shaking dependencies from your app is whether your project's modules have side effects. An example of a side effect is when a function modifies something outside of its own scope, which is a *side effect* of its execution:

```
let fruits = ["apple", "orange", "pear"];

console.log(fruits); // (3) ["apple", "orange", "pear"]

const addFruit = function(fruit) {
  fruits.push(fruit);
};

addFruit("kiwi");

console.log(fruits); // (4) ["apple", "orange", "pear", "kiwi"]
```



In this very basic example, `addFruit` produces a side effect when it modifies the `fruits` array, which is beyond the `addFruit` function's scope.

Side effects also apply to ES6 modules, and that matters in the context of tree shaking. Modules that take predictable inputs and spit out equally predictable outputs without modifying anything outside of their own scope are dependencies we can safely shake if we're not using them. They're self-contained, *modular* pieces of code. Hence, "modules".

Where webpack is concerned, we can hint that a package and its dependencies are free of side effects by specifying `"sideEffects": false` in a project's `package.json` file:

```
{
  "name": "webpack-tree-shaking-example",
  "version": "1.0.0",
  "sideEffects": false
}
```



Alternatively, you can tell webpack which specific files are not side effect-free:

```
{
  "name": "webpack-tree-shaking-example",
  "version": "1.0.0",
  "sideEffects": [
    "./src/utils/utils.js"
  ]
}
```



In the latter example, any file that isn't specified will be assumed to be free of side effects. If you don't want to add this to your `package.json` file, you can also specify this flag in your webpack config via `module.rules`

Importing only what we need

So we told Babel to leave our ES6 modules be, but now we need to make a slight adjustment to our `import` syntax to bring in only the functions we need from the `utils` module. In this guide's example, all we need is `simpleSort`:

```
import { simpleSort } from "../../utils/utils";
```



Using this syntax, we're saying "hey, get me only the `simpleSort` export from the `utils` module." Because we're bringing in only `simpleSort` and not the entire `utils` module into the global scope, we need to change every instance of `utils.simpleSort` to `simpleSort`:

```
if (this.state.sortBy === "model") {  
  json = simpleSort(json, "model", this.state.sortOrder);  
} else if (this.state.sortBy === "type") {  
  json = simpleSort(json, "type", this.state.sortOrder);  
} else {  
  json = simpleSort(json, "manufacturer", this.state.sortOrder);  
}
```



Now that we've done what we need for tree shaking to work, let's step back for a second. This is the webpack output *before* shaking the dependency tree:

Asset	Size	Chunks		Chunk Names
js/vendors.16262743.js	37.1 KiB	0	[emitted]	vendors
js/main.797ebb8b.js	20.8 KiB	1	[emitted]	main



This is the output *after* tree shaking has been put into place:

Asset	Size	Chunks		Chunk Names
js/vendors.45ce9b64.js	36.9 KiB	0	[emitted]	vendors
js/main.559652be.js	8.46 KiB	1	[emitted]	main



While both bundles shrank, it's really the `main` bundle that benefits the most. By shaking off the unused parts of the `utils` module, we've managed to chop about 60% of the code off this bundle. This not only lowers the amount of time the script takes to the download, but processing time as well.

When things aren't so straightforward

In most cases, tree shaking will work in recent versions of webpack if you make these minor changes, but there are always exceptions that can leave you scratching your head. For instance, [Lodash](#) is a bit of a strange case in that tree shaking as it's described in this guide doesn't work. Because of how Lodash is architected, you have to a) install the [lodash-es](#) package in lieu of regular old [lodash](#) and b) use a slightly different syntax (referred to as "cherry-picking") to shake off the other dependencies:

```
// This still pulls in all of lodash even if everything is configured right
import { sortBy } from "lodash";

// This will only pull in the sortBy routine.
import sortBy from "lodash-es/sortBy";
```

If you prefer your `import` syntax to be consistent, you *could* just use the standard `lodash` package, and install [babel-plugin-lodash](#). Once you add the plugin to your Babel config, you can use the typical `import` syntax you would otherwise use to shake unused exports.

If you run into a stubborn library that won't respond to tree shaking, look to see if it exports its methods using the ES6 syntax. If it's exporting stuff in CommonJS format (e.g., `module.exports`), that code won't be tree shakeable by webpack. Some plugins provide tree shaking functionality for CommonJS modules (e.g., [webpack-common-shake](#)), but this may only go so far as [there some CommonJS patterns you just can't shake](#). If you want to reliably shake unused dependencies from your applications, ES6 modules are what you should use going forward.

Go shake some trees!

Whatever mileage you get out of tree shaking depends on your app and its specific dependencies and architecture. Try it! If you know for a fact you haven't set up your module bundler to perform this optimization, there's no harm trying and seeing what the benefit is to your application. Any unused code you can prune from your bundles is a worthwhile optimization.

You may realize a lot of gains from tree shaking, or not very much at all. But by configuring your build system to take advantage of this optimization in production builds and selectively importing only what your application needs, you'll be proactively keeping your applications as thin as possible. That's good for performance, and by extension, your users.

Special thanks to Kristofer Baxter, Jason Miller, Addy Osmani, Jeff Posnick, Sam Saccone, and Philip Walton for their valuable feedback, which significantly improved the quality of this article.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 3.0 License](#), and code samples are licensed under the [Apache 2.0 License](#). For details, see our [Site Policies](#). Java is a registered trademark of Oracle and/or its affiliates.

Last updated July 2, 2018.